

# CSS430

# Process Management

## Textbook Chapter 3

Instructor: Stephen G. Dame  
e-mail: [sdame@uw.edu](mailto:sdame@uw.edu)

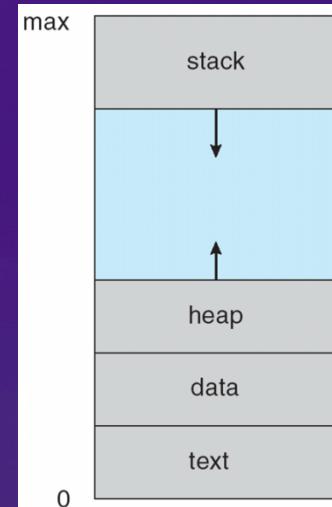
These slides were adapted from the OSC textbook slides (Silberschatz, Galvin, and Gagne),  
Professor Munehiro Fukuda and the instructor's class materials.

“Simple design, **intense** content.”

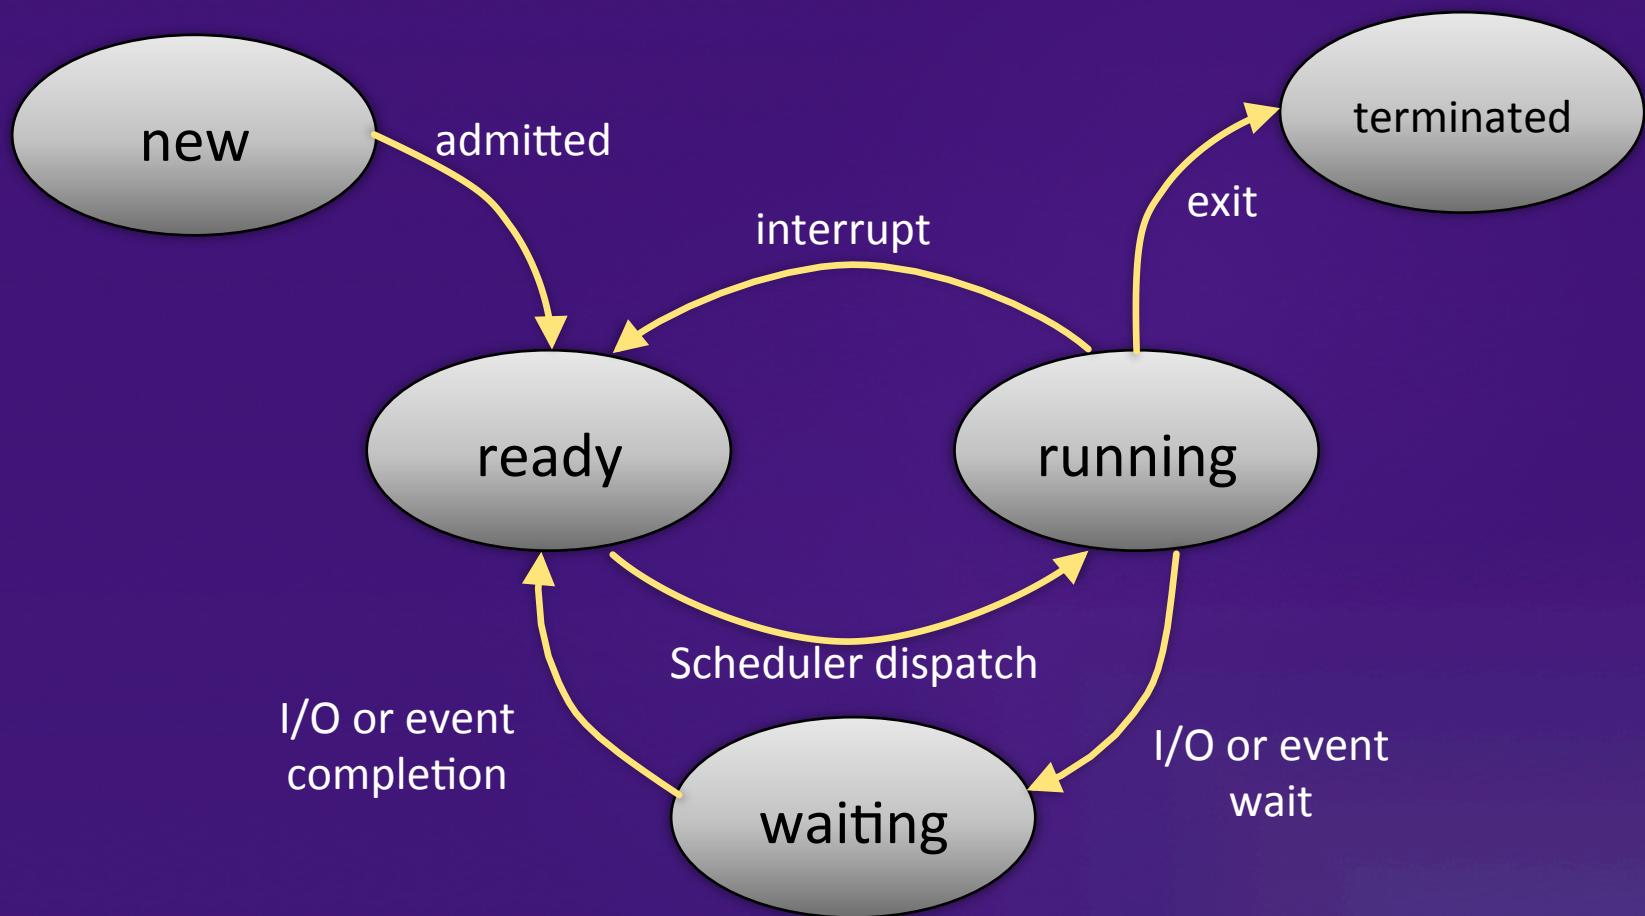
*Edward Rolf Tufte* is an American statistician and professor emeritus of political science, statistics, and computer science at Yale University. He is noted for his writings on information design and as a pioneer in the field of data visualization.

# Process Concept

- ◆ Process – a program in execution; process execution must progress in sequential fashion.
- ◆ Textbook uses the terms ***job*** and ***process*** interchangeably.
- ◆ A process includes:
  - ✓ Program counter
  - ✓ Stack (local variables)
  - ✓ Data section (global data)
  - ✓ Text (code)
  - ✓ Heap (dynamic data)
  - ✓ Files (stdin, stdout, stderr, other file descriptors)

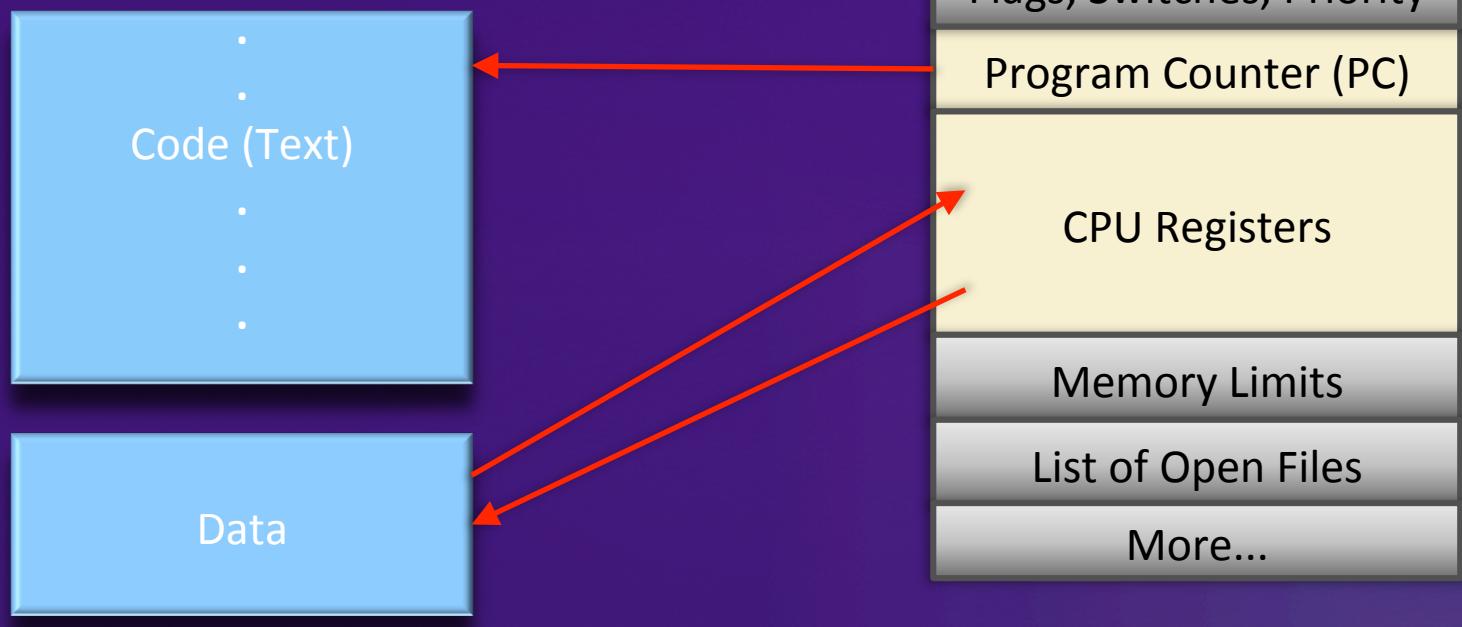


# Process State

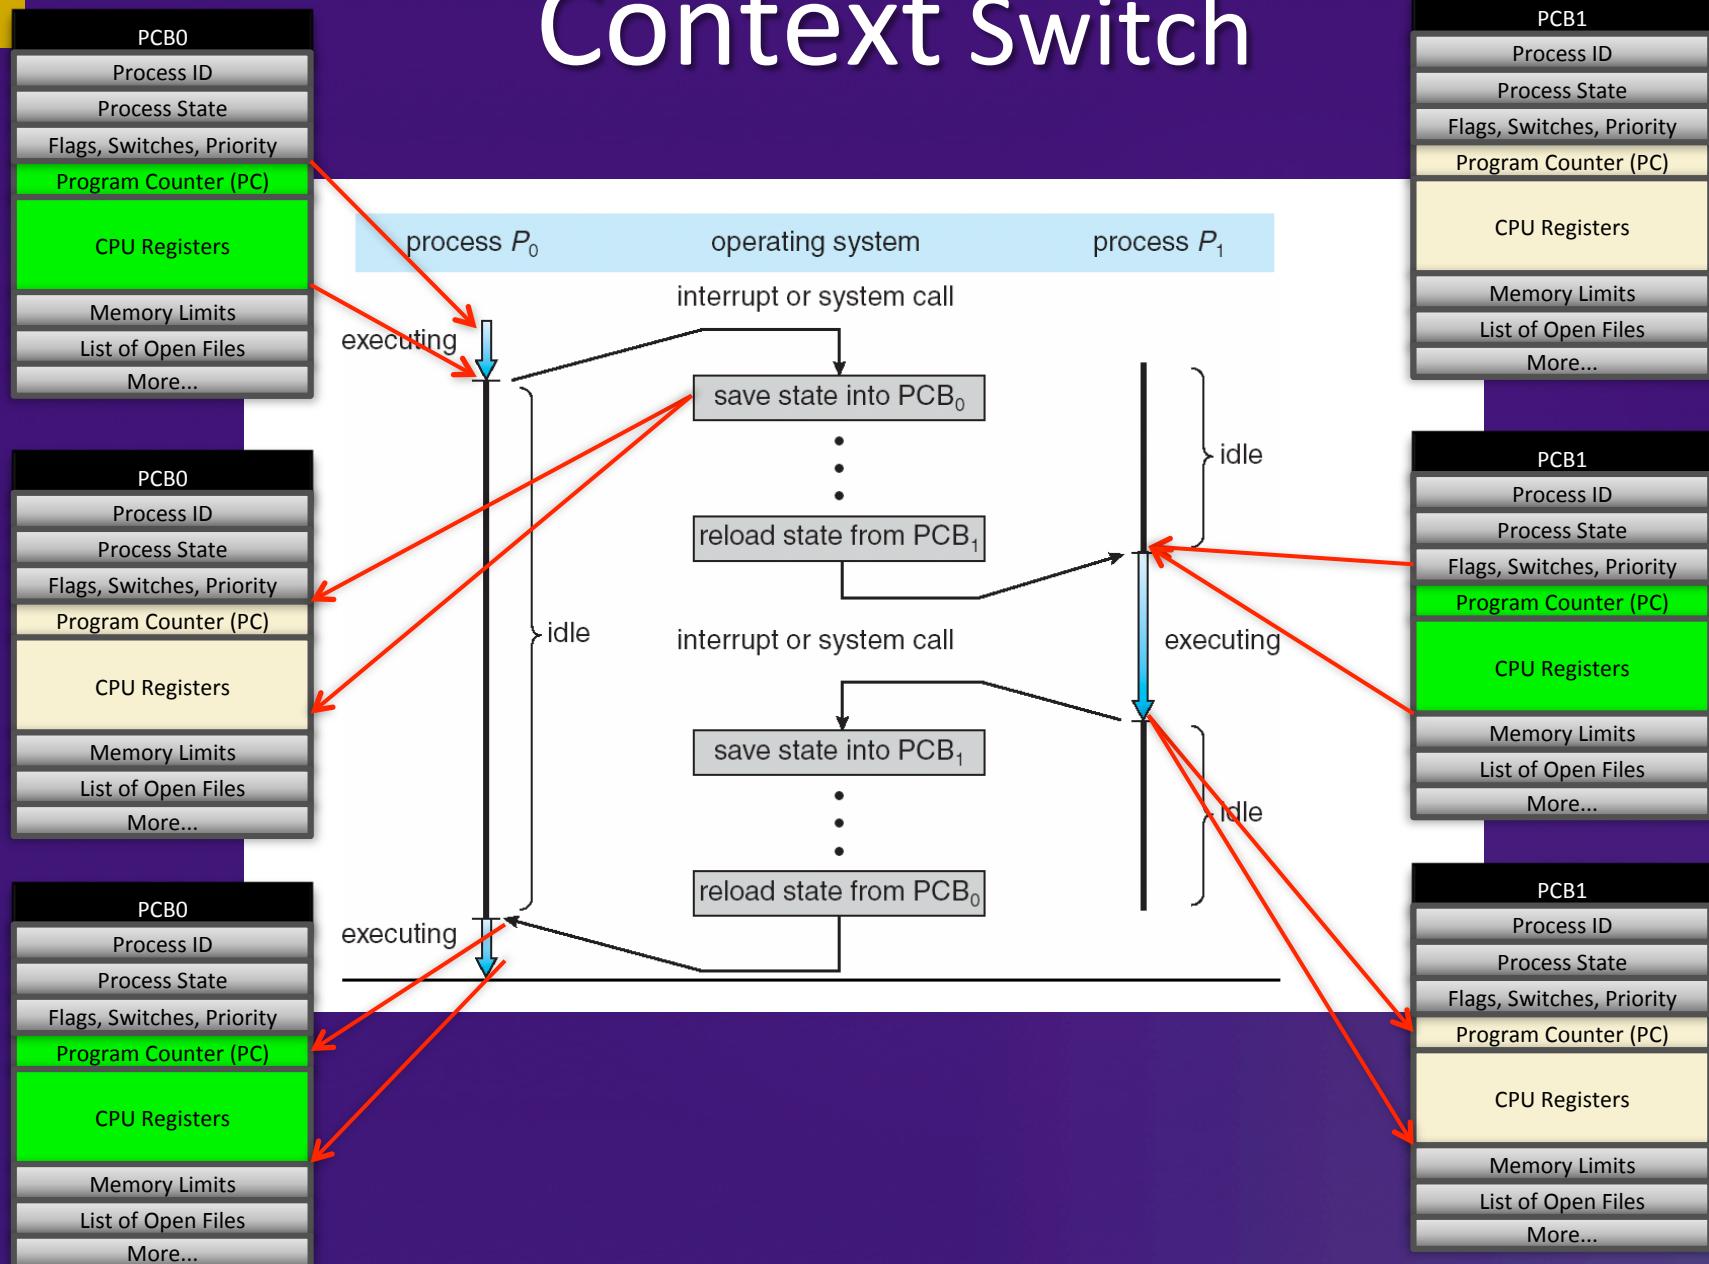


# Process Control Block

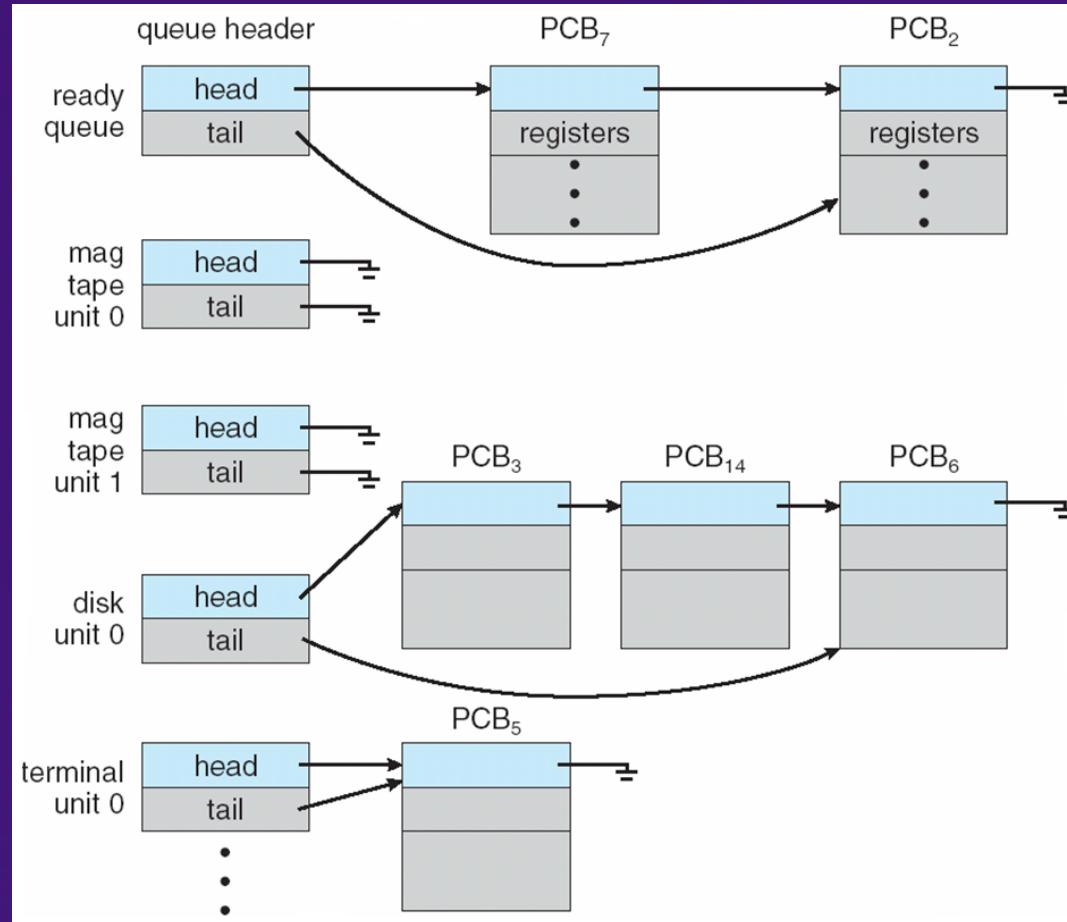
All of the information needed  
to keep track of a process  
when switching context.



# Context Switch



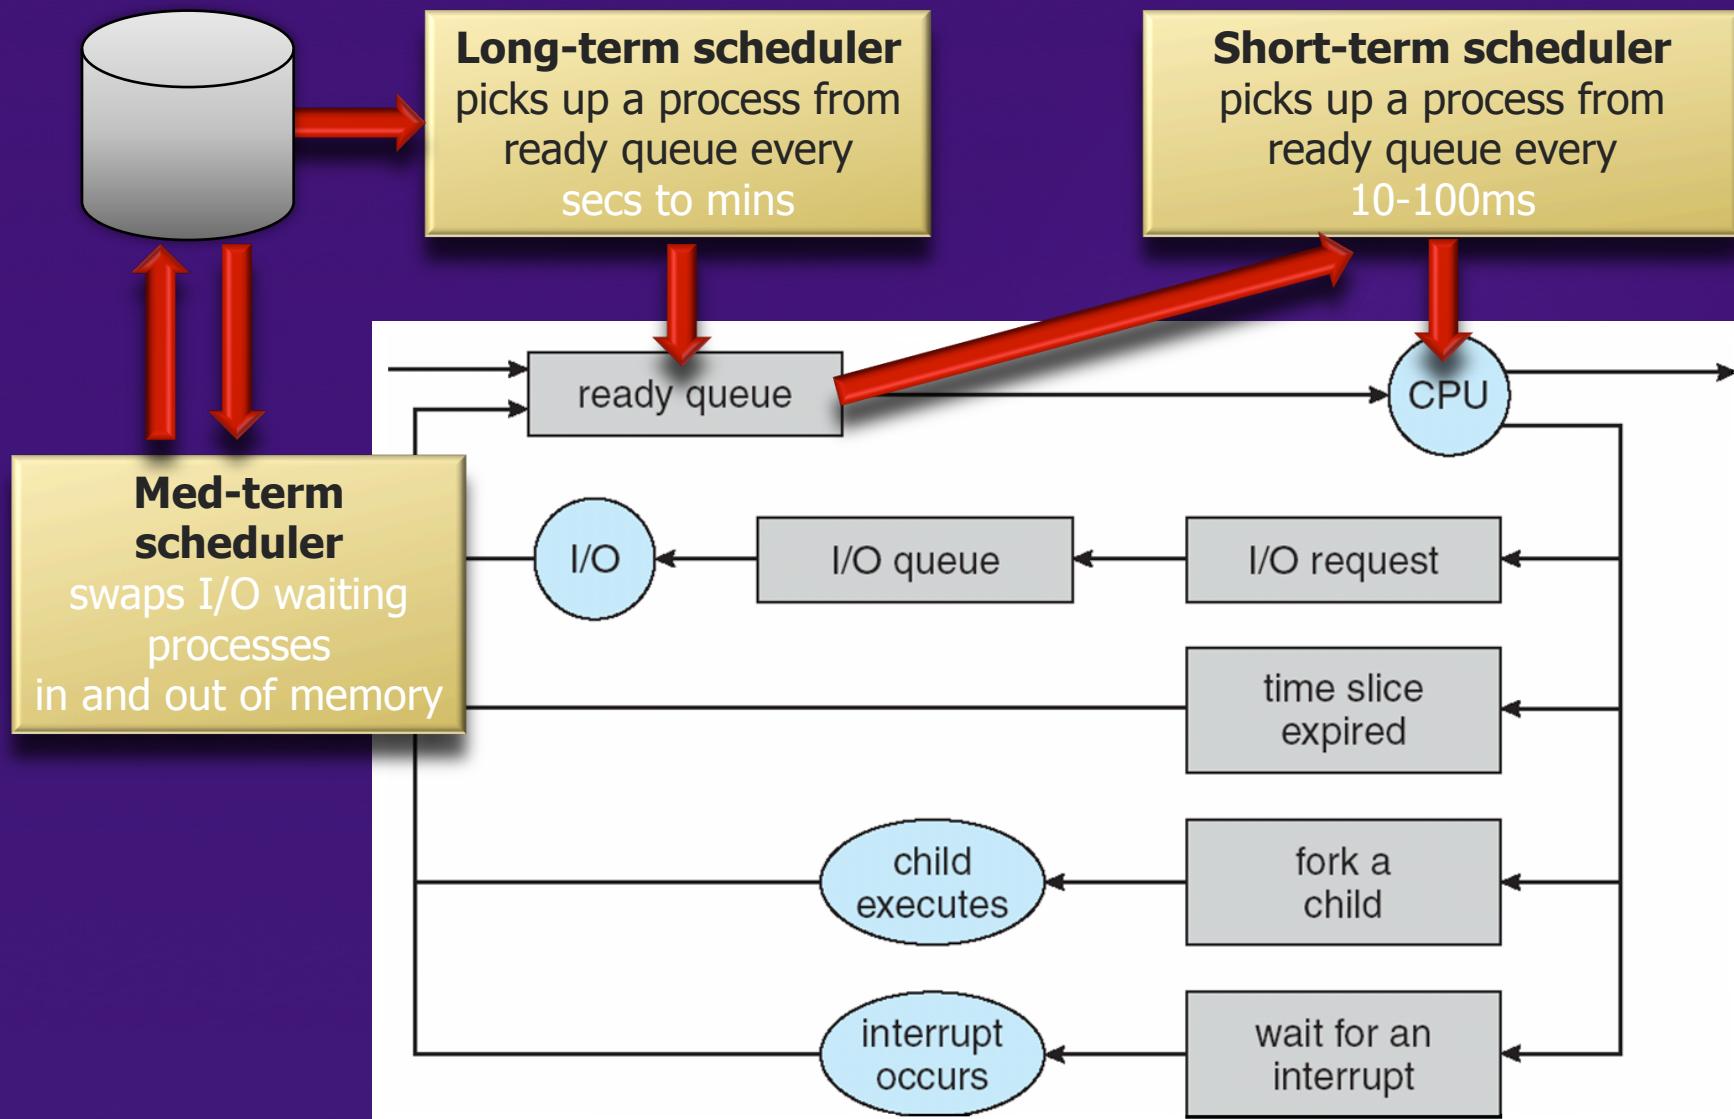
# Process Scheduling Queues



# Process Schedulers

- ◆ Long-Term Scheduler (“Job Scheduler”)
  - ✓ Brings processes into the READY list
  - ✓ Period = seconds to minutes
- ◆ Medium Term Scheduler (“Swapper”)
  - ✓ Swaps inactive processes to disk
  - ✓ Brings back swapped processes on demand
- ◆ Short-Term Scheduler (“CPU Scheduler”)
  - ✓ Selects processes from the READY list
  - ✓ Allocates CPU time to the process
  - ✓ Period = milliseconds

# Representation of Process Scheduling



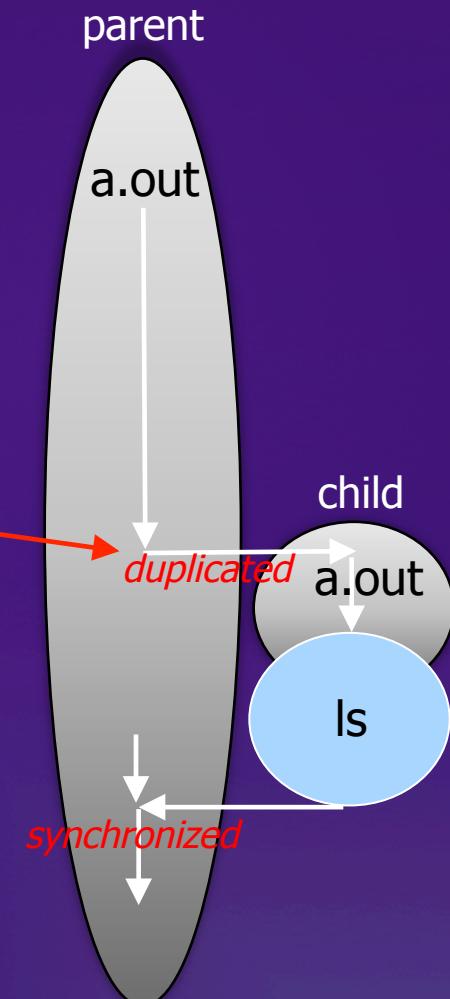
# Process Creation

- ◆ Parent process creates children processes.
- ◆ Resource sharing
  - ✓ Resource inherited by children: file descriptors, shared memory and system queues
  - ✓ Resource **not** inherited by children: address space
- ◆ Execution
  - ✓ Parent and children execute concurrently.
  - ✓ Parent blocks on **wait()** system call until children terminate.
- ◆ UNIX examples
  - ✓ **fork()** system call creates new process.
  - ✓ **execlp** system call used after a **fork()** to replace the process' memory space with a new program.
- ◆ CSS430-unique ThreadOS: SysLib.exec and Syslib.join

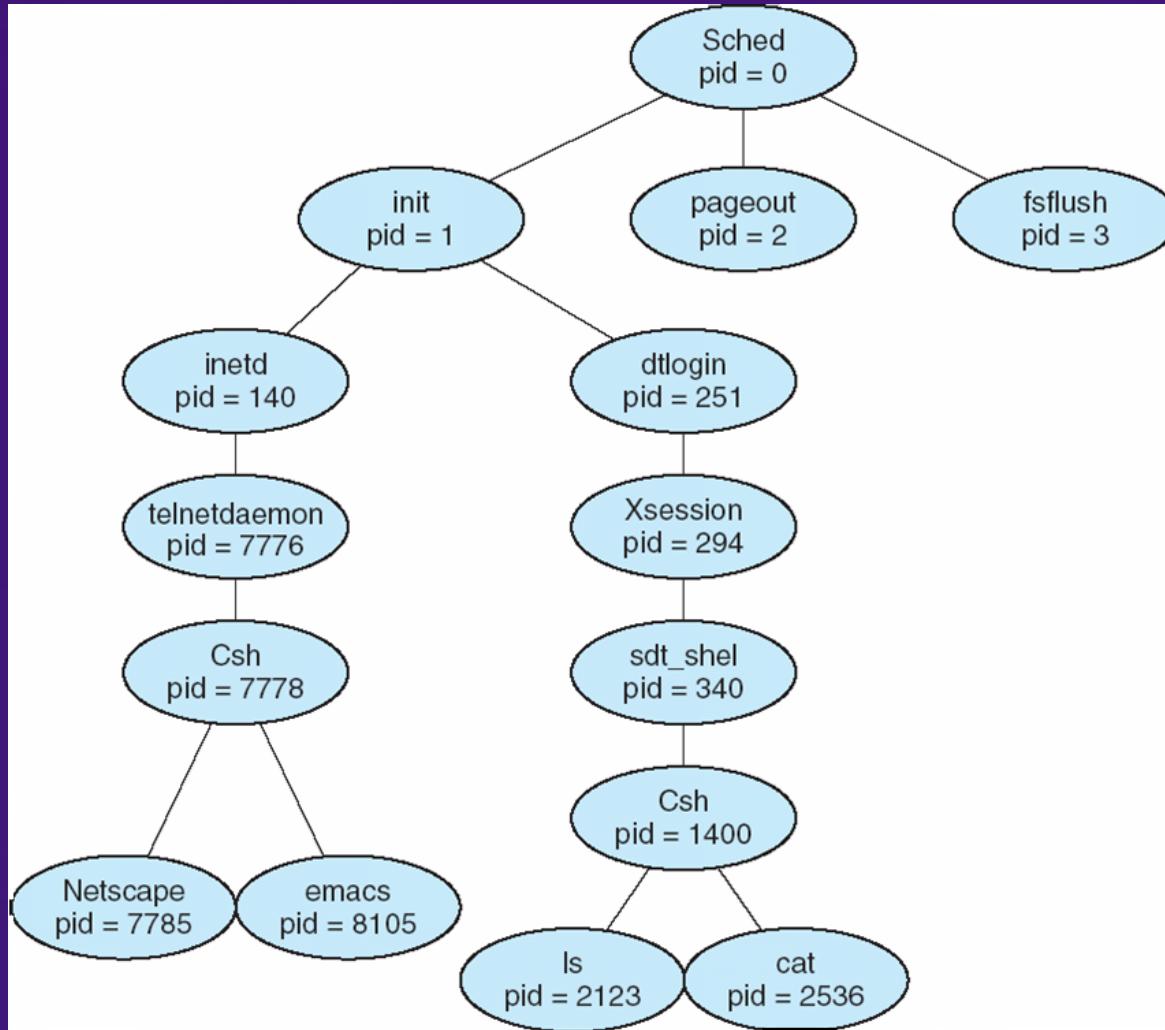
# C - Forking a Separate Process

```
#include <stdio.h> // for printf
#include <stdlib.h> // for exit
#include <unistd.h> // for fork, execvp

int main(int argc, char *argv[])
{
    int pid; // process ID
    // fork another process
    pid = fork();
    if (pid < 0) { // error occurred
        fprintf(stderr, "Fork Failed");
        exit(EXIT_FAILURE);
    }
    // ----- CHILD SECTION -----
    else if (pid == 0) {
        execvp("/bin/ls", "ls", "-l", NULL);
    }
    // ----- PARENT SECTION -----
    else {
        // parent will wait for the child to complete
        wait(NULL);
        printf("Child Complete");
        exit(EXIT_SUCCESS);
    }
}
```



# A Tree of Processes On A Typical Unix System



# Process Termination

- ◆ Process termination occurs when
  - ✓ last statement is executed
  - ✓ **exit()** system called explicitly
- ◆ Upon process termination
  - ✓ Termination code is passed from:  
child (via **exit()**) → parent (via **wait()**).
  - ✓ Process resources are deallocated by OS.
- ◆ Parent may terminate execution of children processes  
(via **kill()**) when:
  - ✓ Child has exceeded allocated resources.
  - ✓ Task assigned to child is no longer required.
  - ✓ Parent is exiting (cascading termination).
    - ❖ Some operating systems do not allow child to continue if its parent terminates.

# Discussion 1

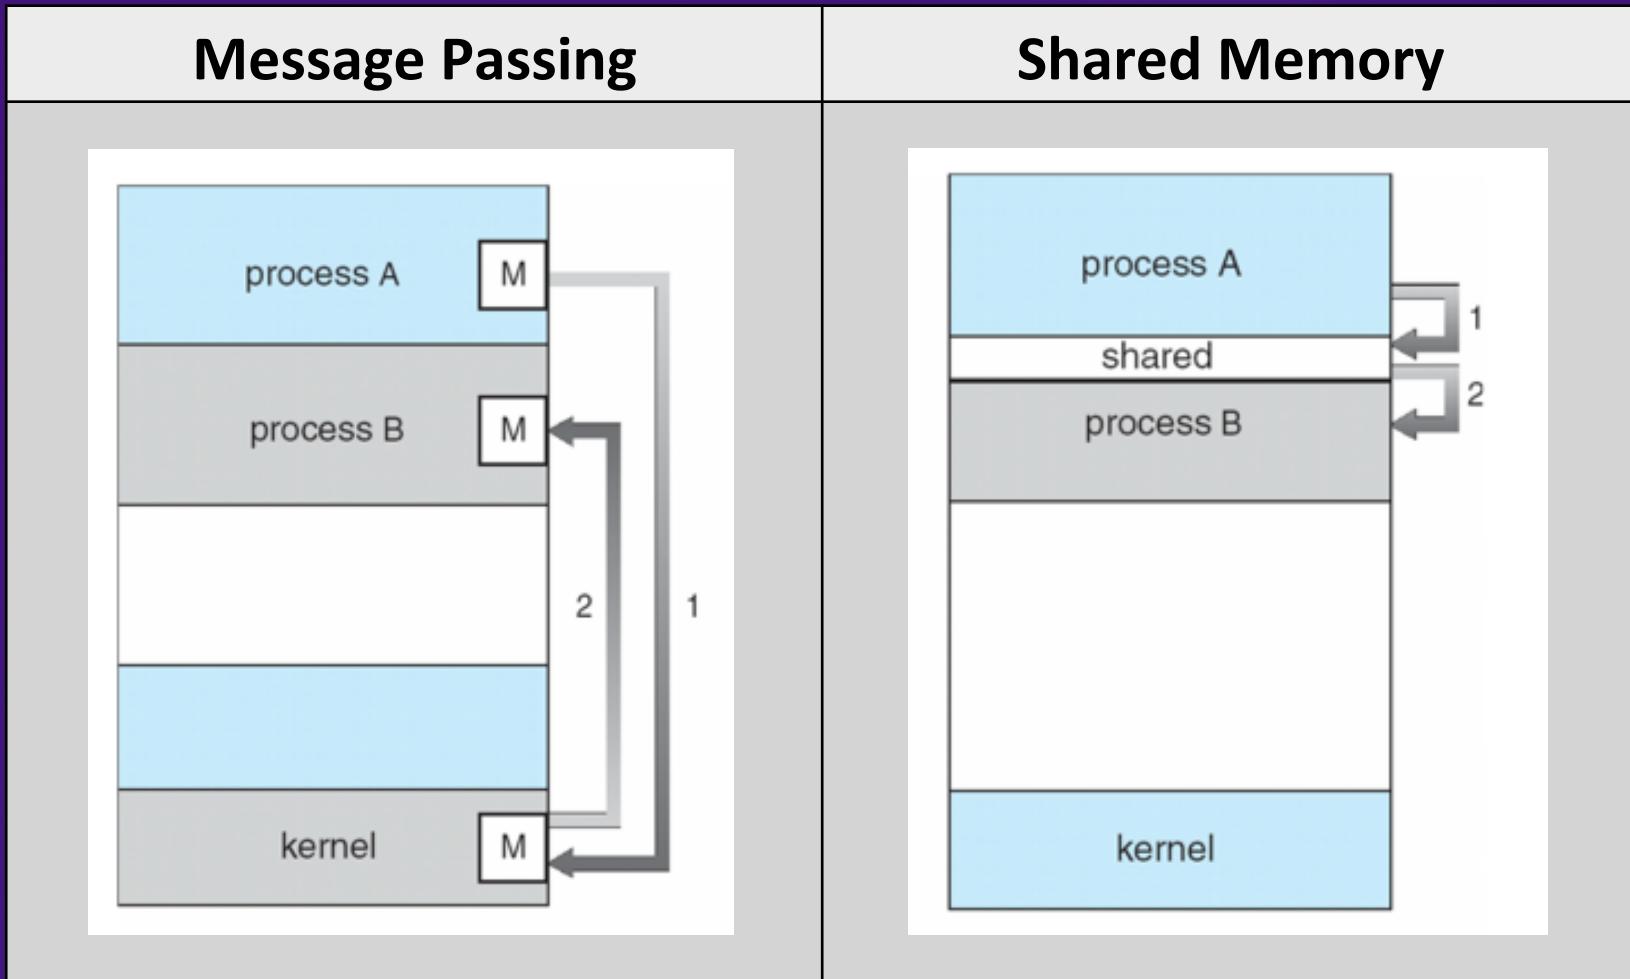
- ① What are the differences between CPU bound and I/O bound processes?
- ② What is another name for the *text section* in a process and what does it contain?
- ③ Discuss the differences between short-term, medium-term, and long-term scheduling in the following table:

Scheduler Type	Source of Processes	(approx.) period of execution (sec)	Key Role
Short Term			
Medium Term			
Long Term			

# Cooperating Processes

- ◆ *Process independency*: Processes belonging to different users do not affect each other unless they give each other certain access permissions
- ◆ *Process Cooperation*: Processes spawned from the same user process share some resources and communicate with each other through these resources (e.g. shared memory, message queues, pipes, and files)
- ◆ Advantages of process cooperation
  - ✓ Information sharing: (sharing files)
  - ✓ Computation speed-up: (parallel programming)
  - ✓ Modularity: (like `who | wc -l`, one process lists current users and another counts the number of users.)
  - ✓ Convenience: (e.g. web-surfing while working on programming with vim and g++)

# Communication Models



# Message Passing

- ◆ Message system – processes communicate with each other without resorting to shared variables.
- ◆ IPC facility provides two operations:
  - ✓ **send**(*message*) – message size fixed or variable
  - ✓ **receive**(*message*)
- ◆ If  $P$  and  $Q$  wish to communicate, they need to:
  - ✓ establish a *communication link* between them
  - ✓ exchange messages via send/receive
- ◆ Implementation of communication link
  - ✓ physical (e.g., shared memory, hardware bus)
  - ✓ logical (e.g., logical properties)

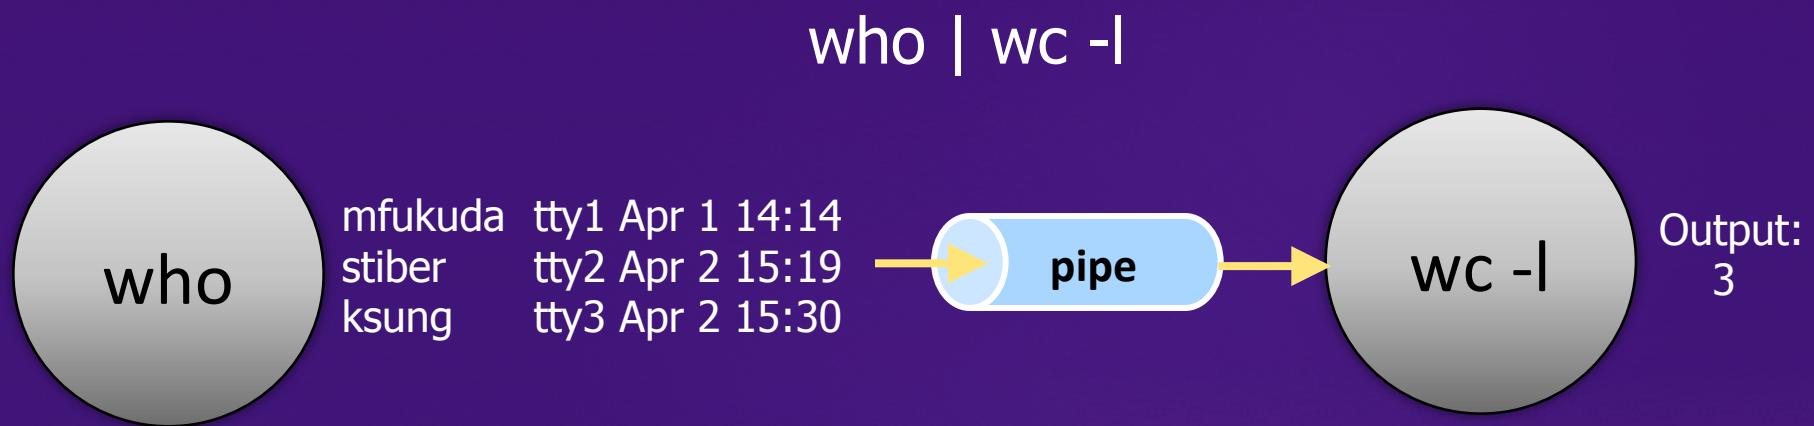
# Direct Communication

# Direct Communication

- ◆ Processes must name each other explicitly:
  - ✓ **send** ( $P, message$ ) – send a message to process P
  - ✓ **receive**( $Q, message$ ) – receive a message from process Q
- ◆ How can a process locate its partner to communicate with?
  - ✓ Processes are created and terminated dynamically and thus a partner process may have gone.
  - ✓ Direct communication takes place between a parent and its child process in many cases.

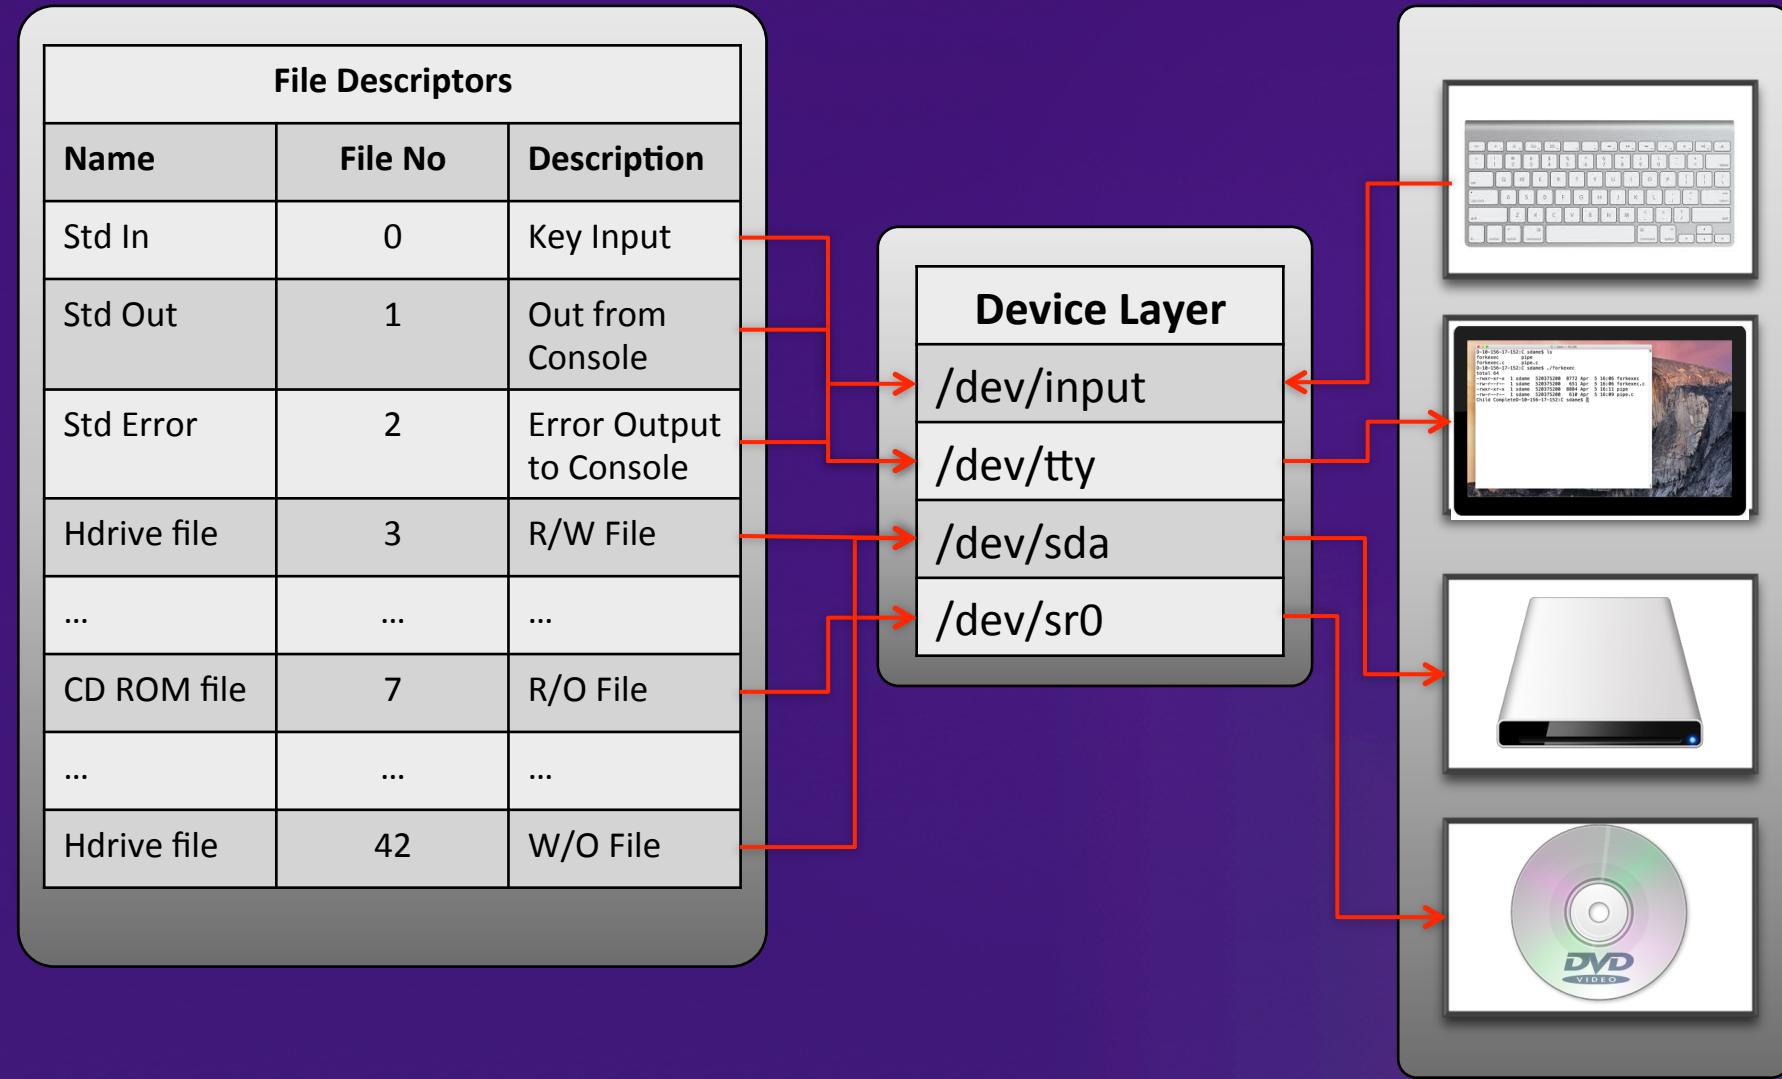
Example: pipe(fd)

# Producer-Consumer Problems



- ◆ Producer process:
  - ✓ **who** produces a list of current users.
- ◆ Consumer process
  - ✓ **wc** receives it for counting # of users (i.e. *lines*).
- ◆ Communication link:
  - ✓ OS provides a pipe.

# File Descriptors



# Direct Communication

Filename: pipe.c

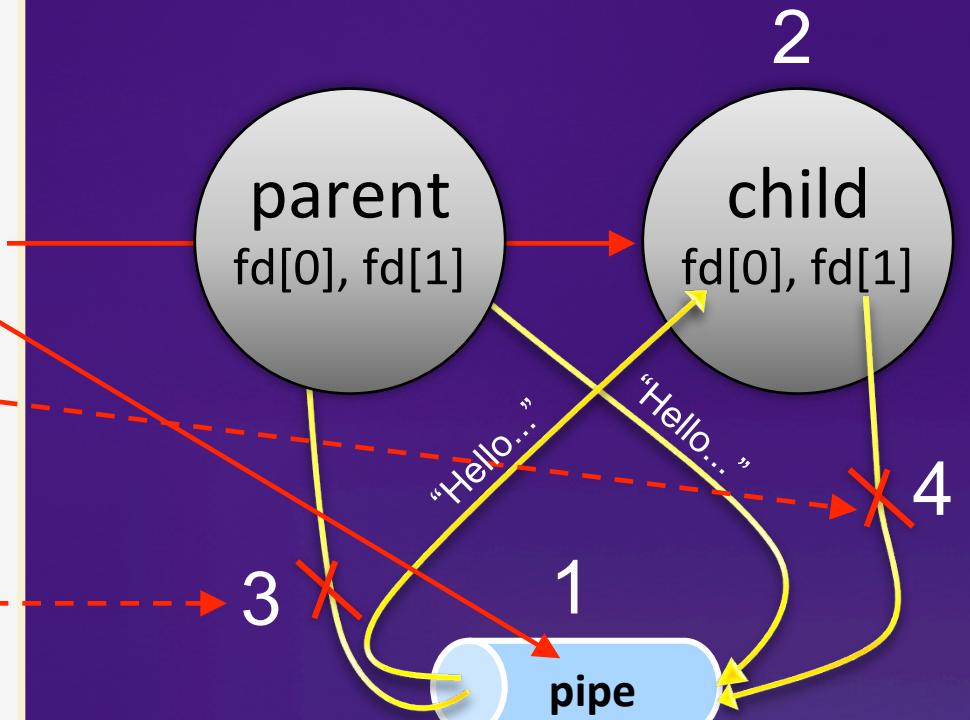
```
#include <unistd.h>      // for fork, pipe

int main( void ) {
    enum {RD, WR}; // pipe fd index RD=0, WR=1
    int n, fd[2];
    pid_t pid;
    char buf[100];

    if( pipe(fd) < 0 ) // 1: pipe created
        perror("pipe error");
    else if ((pid = fork()) < 0) // 2: child forked
        perror("fork error");

    else if (pid == 0) {          Child process
        close(fd[WR]); // 4: child's fd[1] closed
        n = read(fd[RD], buf, 100);
        write(STDOUT_FILENO, buf, n);
    }

    else {                      Parent process
        close(fd[RD]); // 3: parent's fd[0] closed
        write(fd[WR], "Hello my child\n", 15);
        wait(NULL);
    }
}
```



# pipe, dup2 in C++

```
#include <stdlib.h>      // for exit
#include <stdio.h>        // for perror
#include <unistd.h>        // for fork, pipe
#include <sys/wait.h>       // for wait
#include <iostream>         // for cerr, cout
using namespace std;

int main( void ) {
    enum {RD, WR}; // pipe fd index RD=0, WR=1
    int n, fd[2];
    pid_t pid;

    if( pipe(fd) < 0 ) // 1: pipe created
        perror("pipe error");
    else if ((pid = fork()) < 0) // 2: child forked
        perror("fork error");
    else if (pid == 0) {
        close(fd[WR]); // 4: child's fd[1] closed
        dup2(fd[RD], 0); // stdin(0) --> child's pipe read

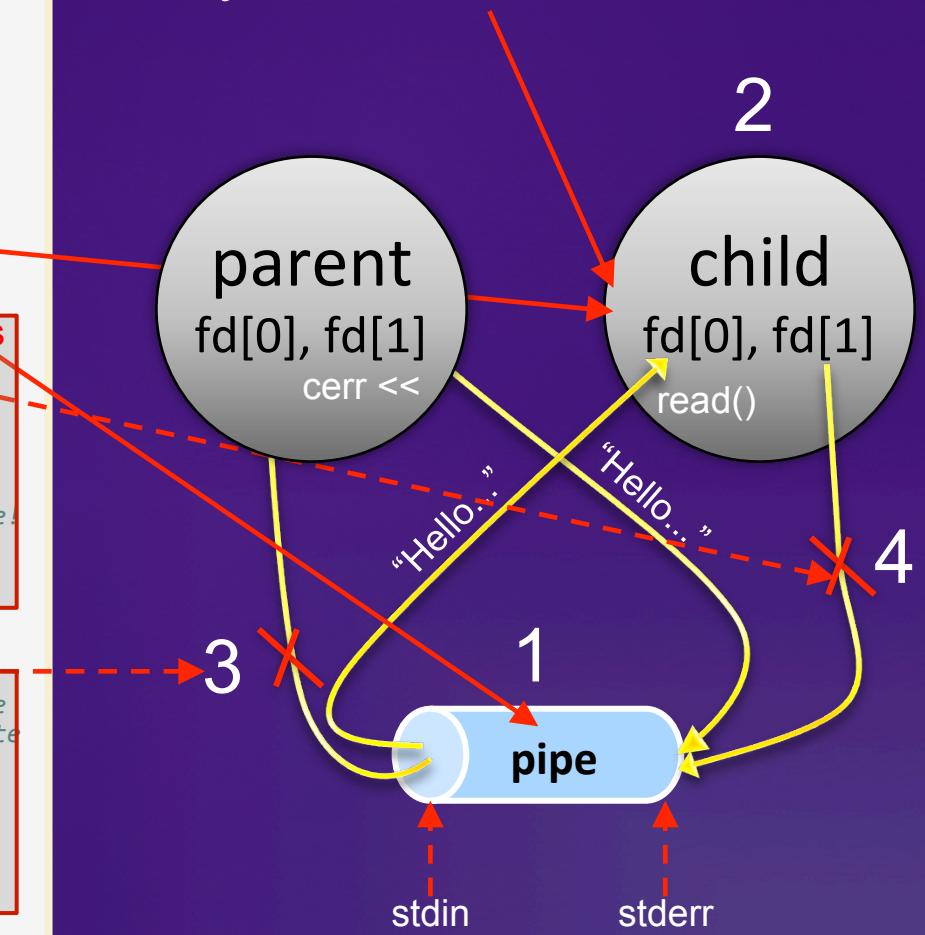
        char buf[256];
        n = read(fd[RD], buf, 256); // use this raw read!
        // cin >> buf; <-- *caution with cin and white space*
        cout << buf;           // write to stdout
        cout << "Child Done!" << endl;
    }
    else {
        close(fd[RD]); // 3: close parent's read end of pipe
        dup2(fd[WR], 2); // stderr(2) --> parent's pipe write

        cerr << "Hello my child" << endl;
        wait( NULL );
        cout << "Parent Done!" << endl;
    }
    exit(EXIT_SUCCESS);
}
```

CSS430 Operating Systems : Process Management

v0.6

Why can't we use `cin >> buf`?



# pipe, dup2, execvp in C++

Filename: pipe\_exec.cpp

```
#include <stdlib.h>      // for exit
#include <stdio.h>        // for perror
#include <unistd.h>       // for fork, pipe
#include <sys/wait.h>      // for wait
#include <iostream>         // for cerr, cout
#define MAXSIZE 4096
using namespace std;

int main( void ) {
    enum {RD, WR}; // pipe fd index RD=0, WR=1
    int n, fd[2];
    pid_t pid;

    if( pipe(fd) < 0 ) // 1: pipe created
        perror("pipe error");
    else {
        switch(fork()) {
            case -1:
                perror("fork error");
            case 0:
                {
                    close(fd[RD]); // 4: child's fd[0] closed
                    dup2(fd[WR], 1); // stdout --> child's pipe write
                    execvp("/bin/ls", "ls", "-l", NULL);
                    // Never returns here!!
                }
        }
    }

    default:
    {
        close(fd[WR]); // 3: close parent's write end of pipe
        dup2(fd[RD], 0); // stdin --> parent's pipe read
        wait( NULL );

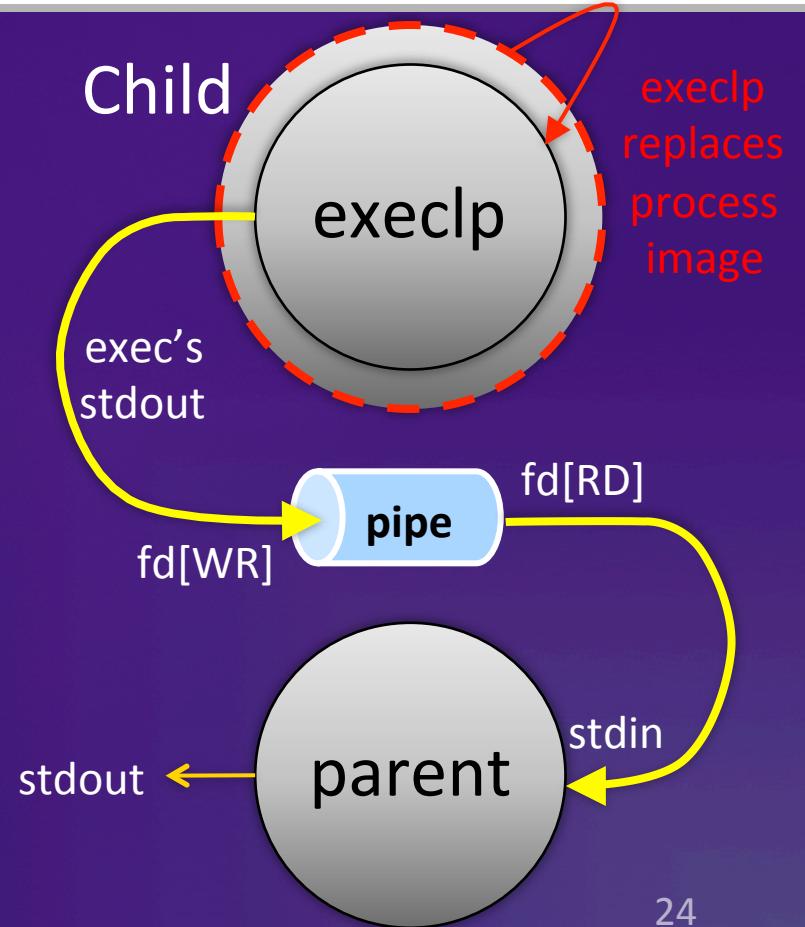
        char buf[MAXSIZE];
        n = read(fd[RD], buf, MAXSIZE); // use this raw read!
        buf[n] = '\0'; // make sure cstring is terminated
        cout << buf; // write to stdout
        cout << "Parent Done!" << endl;
    }
}

exit(EXIT_SUCCESS);
}
```

Child process

Parent process

```
# ./pipe_exec
total 128
-rwxr-xr-x  1 Steve  Steve  9788 Apr  6 18:27 pipe
-rw-r--r--  1 Steve  Steve  1576 Apr  6 17:11 pipe.cpp
-rwxr-xr-x  1 Steve  Steve  9784 Apr  6 18:52 pipe_exec
-rw-r--r--  1 Steve  Steve  1723 Apr  6 18:51
pipe_exec.cpp
Parent Done!
```



# Indirect Communication (non-hard coded IDs)

# Indirect Communication (Message Queues)

- ◆ Messages are directed and received from mailboxes (also referred to as ports).
  - ✓ Each mailbox has a unique id.
  - ✓ Processes can communicate only if they share a mailbox.
- ◆ Processes must know only a **mailbox id**. They do not need to locate their partners
  - ✓ Example: message queue

# Example: Message Queues

msg\_snd.cpp

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#include <iostream>
using namespace std;

struct mymesg {
    long mytype;
    char mtext[512];
} message_body;

int main( void ) {
    int msgid = msgget( 100, IPC_CREAT );
    strcpy( message_body.mtext, "hello world\n" );
    msgsnd( msgid, &message_body, 512, 0 );
}
```

key=100

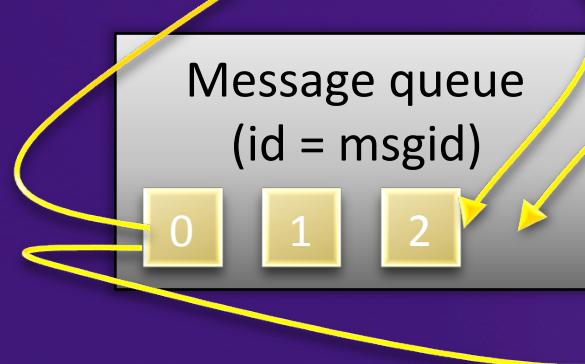
msg\_rcv.cpp

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#include <iostream>
using namespace std;

struct mymesg {
    long mytype;
    char mtext[512];
} message_body;

int main( void ) {
    int msgid = msgget( 100, IPC_CREAT );
    msgrcv( msgid, &message_body, 512, 0, 0 );
    cout << message_body.mtext << endl;
}
```

key=100



Some other process can also enqueue and dequeue a message (on the same key)

# msg\_snd.cpp → msg\_rcv.cpp

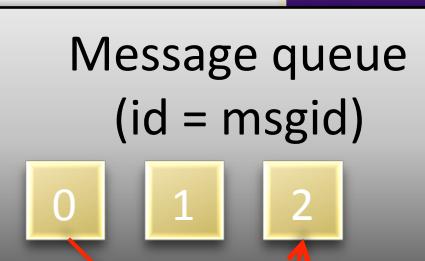
```
#include <stdlib.h> // for exit
#include <stdio.h> // for perror
#include <sys/ipc.h> // for IPC_CREAT
#include <sys/msg.h> // for msgget, msgrcv
#include <iostream> // for cin, cout, cerr
#include <string.h> // for strcpy, strlen
using namespace std;
#define MSGSZ 128

typedef struct {
    long mtype;
    char mtext[MSGSZ];
} message_buf;

main() {
    int msqid;
    size_t msgsz;
    int msgflg = IPC_CREAT | 0666;
    message_buf mymsg;

    key_t key = 74563;
    if((msqid = msgget(key, msgflg)) < 0) {
        perror("msgget");
        exit(1);
    }
    cout << "Sending to Msg Queue:" << key << endl;
    strcpy( mymsg.mtext, "Hello Huskies!" );
    msgsز = strlen(mymsg.mtext) + 1;
    mymsg.mtype = 1;

    if(msgsnd(msqid,&mymsg,msgsز,IPC_NOWAIT) < 0)
    {
        perror("msgsnd");
    }
}
```



```
#include <stdlib.h> // for exit
#include <stdio.h> // for perror
#include <sys/ipc.h> // for IPC_CREAT
#include <sys/msg.h> // for msgget, msgrcv
#include <iostream> // for cin, cout, cerr
using namespace std;
#define MSGSZ 128

typedef struct {
    long mtype;
    char mtext[MSGSZ];
} message_buf;

main() {
    int msqid;
    size_t msgsز;
    int msgflg = IPC_CREAT | 0666;
    message_buf mymsg;

    key_t key = 74563;
    if((msqid = msgget(key, msgflg)) < 0) {
        perror("msgget");
        exit(1);
    }

    if(msgrcv(msqid, &mymsg, MSGSZ, 1, 0) < 0) {
        perror("msgrcv");
        exit(1);
    }
    cout << mymsg.mtext << endl;
}
```

# Shared Mailbox (Java)

```

import java.util.Vector;

public class MessageQueue<E> implements Channel<E>
{
    private Vector<E> queue;

    // a "generic" vector queue
    public MessageQueue() {
        queue = new Vector<E>();
    }
    // nonblocking send
    public void send(E item) {
        queue.addElement(item);
    }
    // nonblocking receive
    public E receive() {
        if (queue.size() == 0)
            return null;
        else
            return queue.remove(0);
    }
}

```

Next chapter will show sharing of this mailbox between threads.

```

public interface Channel<E>
{
    public void send(E item);
    public E receive();
}

```

Makes use of generics <E>  
And Vector<E> to implement a simple Date message mailbox.

```

import java.util.Date;

public class Test
{
    public static void main(String[] args) {
        Channel<Date> mailBox = new MessageQueue<Date>();
        mailBox.send(new Date());

        Date rightNow = mailBox.receive();
        System.out.println(rightNow);
    }
}

```

# Synchronization

- ◆ Sending Process
  - ✓ Non-Blocking – Sends and resumes execution
  - ✓ Blocking – Sender is blocked until message is received or accepted by buffer.
- ◆ Receiving Process
  - ✓ Non-Blocking – receives valid or NULL
  - ✓ Blocking – Waits until message arrives

# Discussion 2

- ① In the previous example, what will happen if you log into uw1-320-lab run msg\_rcv(), then log in from a second shell to uw1-320-lab and run msg\_snd()?
- ② What do Vector<E> and Java Generic do for us?
- ③ What would be two methods of connecting machines together geographically separated from an indirect ID perspective?

# Buffering\*\*

- ◆ Established by creating a **shared memory** area between a *Producer* and a *Consumer*.
- ◆ Queue of messages attached to the link; implemented in one of three ways:
  1. **Zero** capacity – 0 messages  
Producer must wait for Consumer (rendezvous).
  2. **Bounded** capacity – finite length of  $n$  messages  
Producer must wait if link is full (This happens in practical world like sockets).
  3. **Unbounded** capacity – “infinite” length  
Producer never waits. (Non-blocking send)

\*\* Applies to both Direct and Indirect Communications

# Shared Memory

## *Bounded Buffer*

- ◆ Communication link or media has a bounded space to buffer in-transfer data.

```
import java.util.*;

public class BoundedBuffer
{
    public static final int BUFFER_SIZE = 5;
    private int count;      // #items in buffer
    private int in;         // points to the next free position
    private int out;        // points to the next full position
    private Object[] buffer; // a reference to buffer

    public BoundedBuffer( ) {
        count = 0;
        in = 0;
        out = 0;
        buffer = new Object[BUFFER_SIZE];
    }
    public void enter( Object item ) {...}
    public object remove( ) {...}
}
```

See also (for C++ references):

[Beej's Guide to Unix IPC - shm](#)

# Producer and Consumer Processes

## Producer Process

```
for(int i = 0; i++ ) {  
    BoundedBuffer.enter(new Integer(i));  
}
```

```
public void enter( Object item ) {  
    while ( count == BUFFER_SIZE )  
        ; // buffer is full! Wait till buffer is consumed  
    ++count;  
    buffer[in] = item; // add an item  
    in = ( in + 1 ) % BUFFER_SIZE; // circular buffer  
}
```

## Consumer Process

```
for(int i = 0; ; i++ ) {  
    BoundedBuffer.remove();  
}
```

```
public object remove( ) {  
    Object item;  
    while ( count == 0 )  
        ; // buffer is empty! Wait till buffer is filled  
    -- count;  
    item = buffer[out]; // pick up an item  
    out = ( out + 1 ) % BUFFER_SIZE; // circular buffer  
}
```

Buffer[0] [1] [2] [3] [4]



out=1

in=4

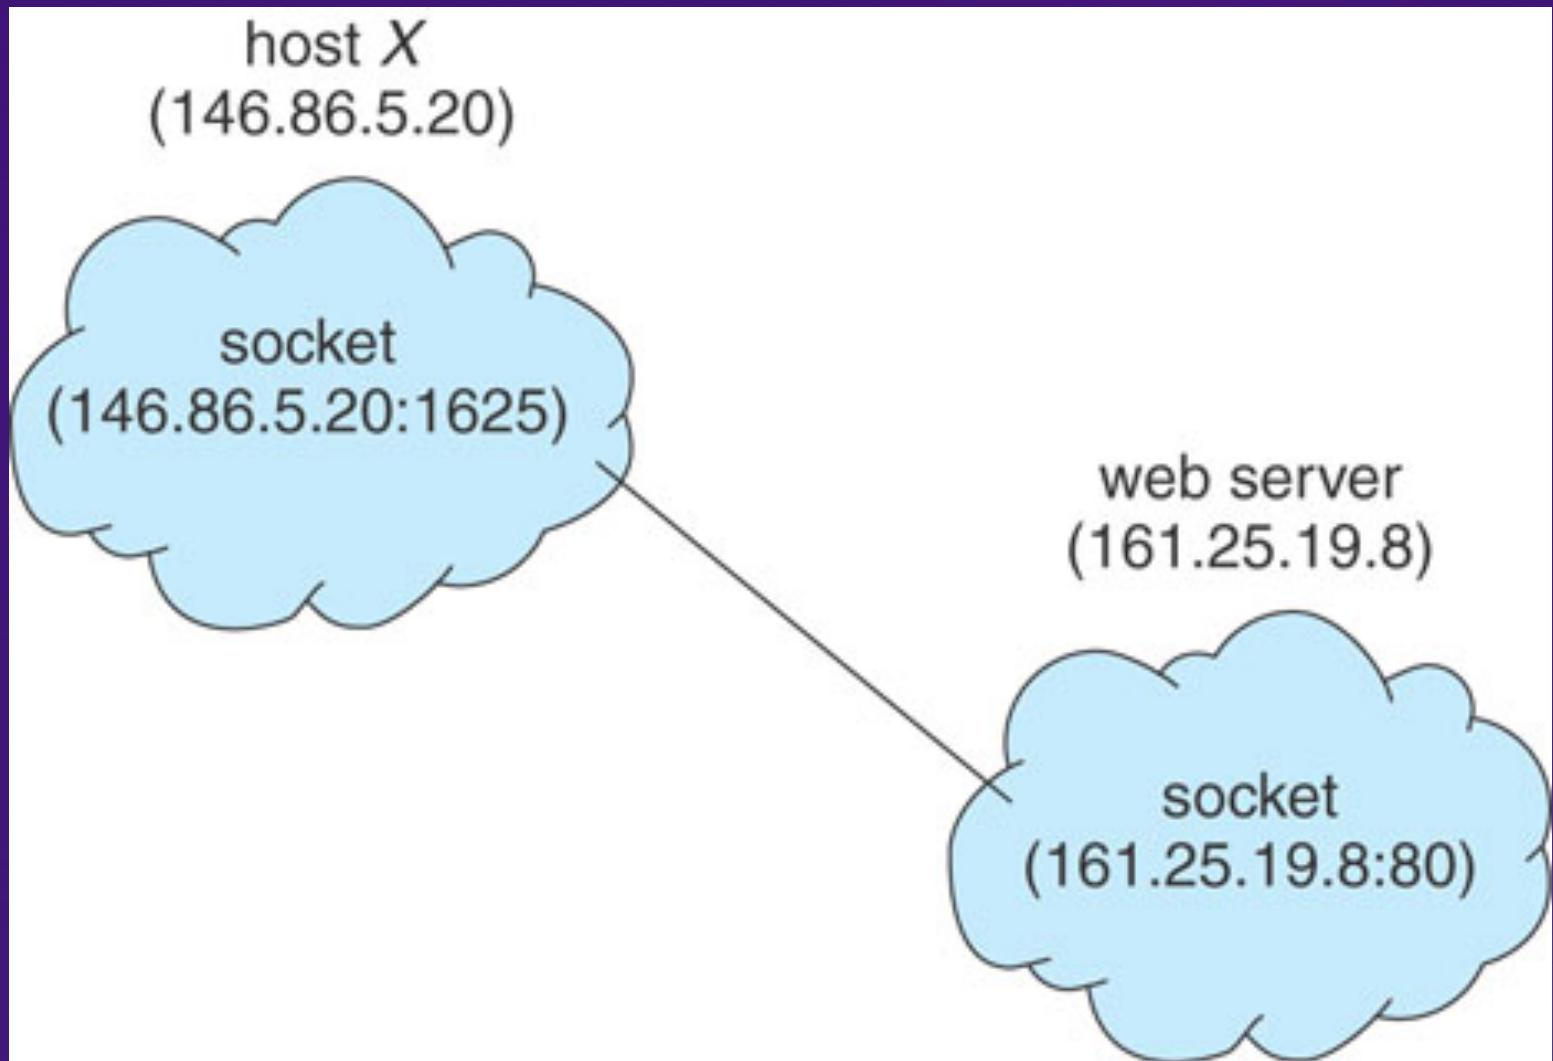
# Client-Server Systems

- ◆ Sockets
- ◆ Remote Procedure Calls (RPC)
- ◆ Remote Method Invocation (RMI)

# Sockets

- ◆ Common way to connect hosts at different locations and/or IP
- ◆ Low-level form of communication between distributed processes
- ◆ Allows only an unstructured stream of data

# Socket Connection



# Datagram Sockets

- ◆ User Datagram Protocol (UDP)
- ◆ Also know as “connectionless” sockets
- ◆ (can be) Point to multi-point
- ◆ Not guaranteed to be delivered
- ◆ Java Classes for UDP:  
DatagramPacket, DatagramSocket,  
MulticastSocket

# Stream Sockets

- ◆ Transmission Control Protocol (TCP)
- ◆ Also known as “connection-oriented” sockets
- ◆ Point to point
- ◆ Guaranteed to be delivered
- ◆ (ftp, http, telnet, ssh, etc.)
- ◆ Java Classes for TCP:  
URL, URLConnection, Socket, SocketServer

# (TCP) Socket Client-Server

```
import java.net.*;
import java.io.*;
public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);
            // now listen for connections
            while (true) {
                Socket client = sock.accept();
                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);
                // write the Date to the socket
                pout.println(new java.util.Date().toString());
                // close the socket and resume
                // listening for connections
                client.close();
            }
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

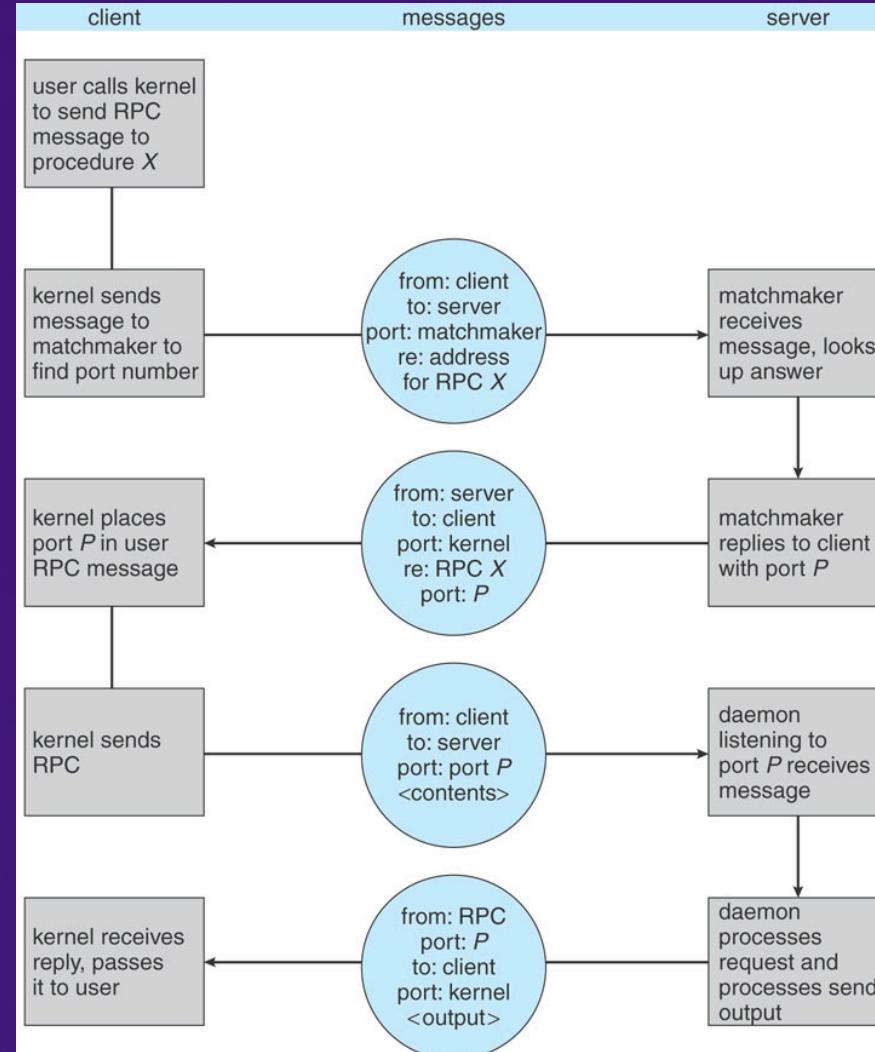
Server

```
import java.net.*;
import java.io.*;
public class DateClient
{
    public static void main(String[] args) {
        try {
            //make connection to server socket
            Socket sock = new Socket("127.0.0.1",6013);
            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));
            // read the date from the socket
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);
            // close the socket connection
            sock.close();
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

Client

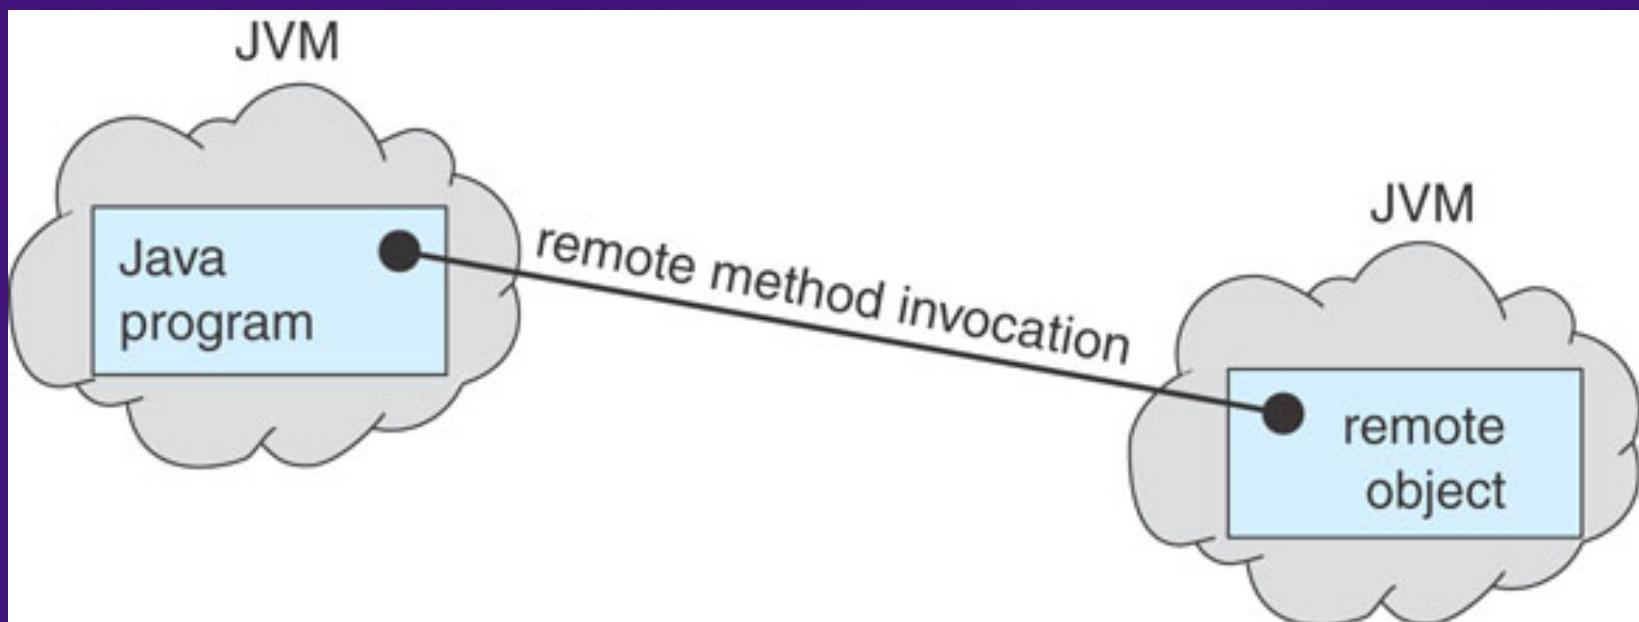
# RPC

- C/C++ Only
- UDP Based

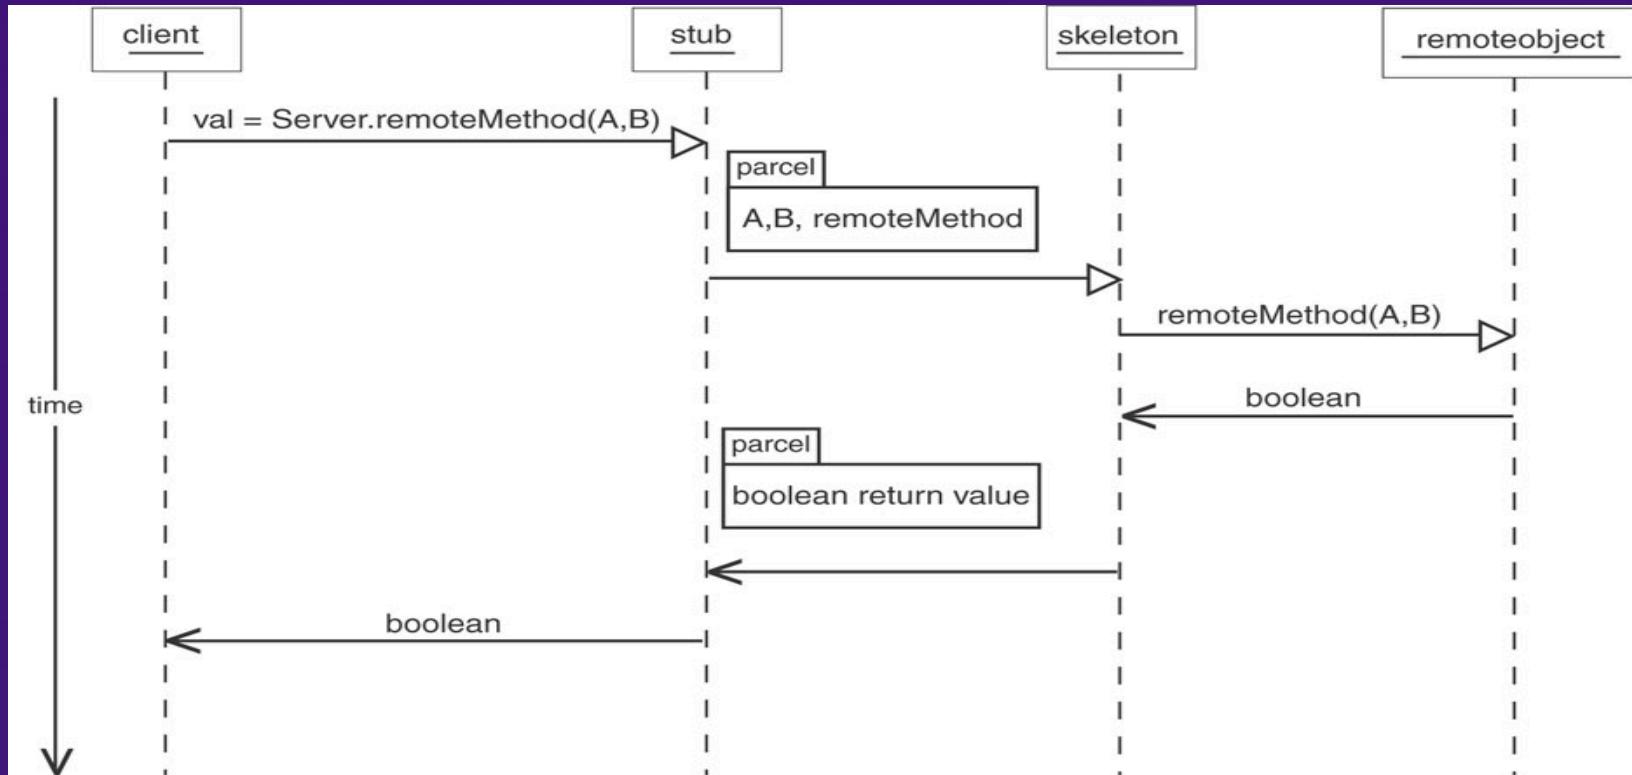


# Remote Method Invocation (RMI)

- Java Based
- TCP Based Communications (i.e. more reliable)



# Sequence Diagram (RMI)



- ◆ Stub provides an interface on the client side, which maps to a port on the remote side.
- ◆ Marshalling - involves packaging the parameters into a form that can be transmitted over a network

# ThreadOS SysLib

```
public interface SysLib {
    public static int exec( String args[] );
    public static int join( );
    public static int boot( );
    public static int exit( );
    public static int sleep( int milliseconds );
    public static int disk( );
    public static int cin( StringBuffer s );
    public static int cout( String s );
    public static int cerr( String s );
    public static int rawread( int blkNumber, byte[] b );
    public static int rawwrite( int blkNumber, byte[] b );
    public static int sync( );
    public static int cred( int blkNumber, byte[] b );
    public static int cwrite( int blkNumber, byte[] b );
    public static int flush( );
    public static int csync( );
    public static String[] stringToArgs( String s );
    public static void short2bytes( short s, byte[] b, int offset );
    public static short bytes2short( byte[] b, int offset );
    public static void int2bytes( int i, byte[] b, int offset );
    public static int bytes2int( byte[] b, int offset );
}
```

# Shell.java (starter example)

```
import java.io.*;
import java.util.*;

class Shell extends Thread
{
    //command line string to contain the full command
    private String cmdLine;

    // constructor for shell
    public Shell( ) {
        cmdLine = "";
    }

    // required run method for this Shell Thread
    public void run( ) {

        // build a simple command that invokes PingPong
        cmdLine = "PingPong abc 100";

        // must have an array of arguments to pass to exec()
        String[] args = SysLib.stringToArgs(cmdLine);
        SysLib.out("Testing PingPong\n");

        // run the command
        int tid = SysLib.exec( args );
        SysLib.out("Started Thread tid=" + tid + "\n");

        // wait for completion then exit back to ThreadOS
        SysLib.join();
        SysLib.out("Done!\n");
        SysLib.exit();
    }
}
```

# Week 2 Homework

1. In Unix, the first process is called **init**. All the others are descendants of “init”. The **init** process spawns a **sshd** process that detects a new secure ssh requested connection (WKPort 22). Upon a new connection, **sshd** spawns a login process that then loads a **shell** on it when a user successfully logs into the system. Now, assume that the user types **who | grep <username> | wc -l**. Draw a process tree from **init** to those three commands. Add **fork**, **exec**, **wait**, and **pipe** system calls between any two processes affecting each other.
2. Consider four different types of inter-process communication (IPC).
  - a) Pipe: implemented with pipe, read, and write
  - b) Socket: implemented with socket, read, and write
  - c) Shared memory: implemented shmget, shmat, and memory read/write
  - d) Shared message queue: implemented with msgget, msgsnd, and msgrcv
  1. Which types are based on direct communication?
  2. Which types of communication do not require parent/child process relationship?
  3. If we code a produce/consumer program, which types of communication require us to implement process synchronization?
  4. Which types of communication can be used to communicate with a process running on a remote computers?
  5. Which types of communication must use file descriptors?
  6. Which types of communication need a specific data structure when transferring data?